

**IMPORTANT: For all projects, any dataset you use or generate must be uploaded to a shared drive.**

This includes:

- GPS datasets
- Synthetic datasets you generated
- Any additional datasets collected from the internet
- Labeled PGN chess games

You may use the university provided drive (up to 2 TB).

If you have Gemini Pro, you may also use your university Google Drive.

---

## GPS Dataset Format

The dataset must follow this structure:

`dataset_root/`

  └── `images/`

  └── `gt.csv`

The `gt.csv` file must contain (3 columns):

1. `image_name.png / image_name.jpg`
2. Latitude
3. Longitude

Each row must correspond to one image in the `images/` folder.

---

# Final Report Guidelines (All Projects)

Each project submission must include a **final written report** describing the problem, methodology, experiments, and conclusions.

The report should be written in the **style of a scientific paper**, following the structure outlined below.

---

## Length & Format

- **Maximum length:** up to **25 pages**
- **Font size:** 12 pt (or standard LaTeX article/conference defaults)
- **Format:** PDF
- **Language:** English

### Important:

You are **not required** to use all 25 pages.

Do **not** artificially inflate the report length.

Extra pages are intended for **figures, tables, experiments, and ablations**.

---

## Recommended Report Structure

### 1. Abstract

- Concise summary of:
    - The problem
    - Your approach
    - Main results
- 

### 2. Introduction

- Task description and motivation
  - Challenges and goals
  - Main contributions
- 

### 3. Related Work

- Relevant papers, datasets, repositories

- Academic papers (if you used it)
  - GitHub repositories
  - Dataset websites
  - Clarify differences from prior work if you built upon that.
- 

## 4. Method

- Model architecture
  - Input/output representation
  - Training procedure
  - Loss functions
  - Preprocessing/postprocessing
  - Diagrams encouraged
- 

## 5. Experiments

- Dataset description and splits
  - Evaluation metrics
  - Baselines and comparisons
  - Quantitative results (tables, plots, numbers)
  - Qualitative results (visualizations)
- 

## 6. Ablation Study (Required)

If your method consists of components (e.g., **X**, **Y**, **Z**), you **must** justify them via ablation.

- Remove **Y** → what happens?
- Remove **Z** → what happens?
- Provide:
  - Comparison tables
  - Clear explanations in text

The goal is to demonstrate that each component is **beneficial and necessary**.

---

## 7. What Did Not Work (Optional)

You are encouraged to include a short section describing:

- Ideas you tried that **did not work**
- Approaches that failed or underperformed
- Design choices you abandoned and why

This section may include:

- Negative results
- Failed experiments
- Insights learned from mistakes

**Important:**

You are **not penalized** for things that did not work.

On the contrary, clear and honest analysis of failures is viewed positively and reflects good research practice.

---

## 8. Discussion / Limitations (Optional but Recommended)

- Failure cases
  - Limitations of the method
  - Possible future improvements (if you had more time)
- 

## 9. References

- Papers, datasets, repositories, websites
- 

## Code Submission (Required)

A GitHub repository must be provided, containing:

- All source code
- `requirements.txt`
- Clear instructions for:
  - Environment setup (from git clone to installations)
  - Training
  - Running inference / evaluation

A `README.md` must explain how to reproduce results.

---

# Project 4 – Image-to-GPS Regression

---

## Overview

In Project 4, each submission must implement a final evaluation function that receives a **single RGB image** and returns a **GPS coordinate prediction: latitude and longitude**.

This function will be used for **automatic evaluation**, therefore its **signature, input format, output format, and conventions must be followed exactly**.

---

## Required Function

Students must implement the following function:

```
def predict_gps(image: np.ndarray) -> np.ndarray:  
    """  
    Predict GPS latitude and longitude from a single RGB image.  
    """
```

---

## Input Specification

- **image**
  - Type: `numpy.ndarray`
  - Shape: `(H, W, 3)`
  - Channel order: **RGB** (NOT BGR)
  - Dtype: `uint8`
  - Value range: `[0, 255]`

The image corresponds to a real location from the dataset. Images may include viewpoint changes, illumination differences, occlusions, and motion blur.

---

# Output Specification

The function must return a **NumPy array** containing the predicted GPS coordinate.

## Array Requirements

- Type: `numpy.ndarray`
  - Shape: `(2, )`
  - Dtype: floating type (`float32`)
  - Format:
    - `output[0] = latitude`
    - `output[1] = longitude`
- 

## Important Notes

- The returned values must be the **absolute GPS coordinates** (e.g, not relative meters).
  - Predictions must be **as accurate as possible** because evaluation is done against the **exact ground-truth GPS** for each images. (which means “5 meter” error is considered **perfect (and probably not achievable)** results, as GPS give accurate location result up to 5 meters. I will take some margin (some more mistake than 5 meters will still be considered as **perfect result**, but I will not provide you with that number!). I will do that since different phones have different focal lengths and field of views.
- 

## Example Output

```
np.array([31.262345, 34.803210], dtype=np.float32)
```

---

## Additional Constraints

- ✗ Do not return a list/tuple (must be `np.ndarray`)
  - ✗ Do not return strings
  - ✗ Do not return meters or any relative coordinate system
  - ✓ No interactive input
  - ✓ Must run end-to-end from a single function call
- 

## Evaluation Notes

The evaluation system will:

- Call `predict_gps(image)`
- Verify output:
  - Type is `np.ndarray`
  - Shape is `(2, )`
  - Values are finite floats
  - Latitude/longitude are within valid ranges
- Compare against the ground truth GPS for each test image (exact latitude/longitude format from the dataset).

Any deviation from the specification may result in **automatic failure**.